



06. Mechanism: Limited Direct Execution

How to efficiently virtualize the CPU with control?

Issue

Performance (성능)

Control (제어 문제)

Direct Execution

띠용용?

3번 단계) Program을 Memory에 로드

4번 단계) Stack Set-up

Recap: Process Creation

1. Load

2. Run-Time Stack Set-Up

3. Heap Creation

4. 추가적인 initialization tasks

5. Start the program

Issues

Problem 1 : Process가 Restricted Operations를 수행하고자 할 때 문제가 발생

Solution 1: Grant everything

Solution 2: Use protected control transfer

System Call

Register syscall handler at IDT

Interrupt descriptor table

Call flow(Pintos): write()

Key Concept of System Call

Trap instruction

Return-from-Trap instruction

What code to run in trap?

Sources of trap

Trap handler

Trap table

Limited Direction Execution

Problem 2: Switching Between Processes

Solution 1: A cooperative Approach: Wait for system calls

Solution 2: A Non-Cooperative Approach: OS Takes Control

A timer interrupt

Saving and Restoring Context

Context Switch

Limited Direct Execution Protocol (Timer interrupt)

Worried About Concurrency?

How to efficiently virtualize the CPU with control?

OS는 physical CPU를 **time sharing**를 통해 공유해야함

Issue

Performance (성능)

시스템에 과도한 overhead(부담)을 주지 않으면서 어떻게 virtualization을 구현할까?

Control (제어 문제)

CPU에 대한 제어권(control)을 유지하면서 어떻게 프로세스를 효율적으로 실행할 수 있을까?

- 제어권이 없다면 프로세스가 영원히 실행되어 시스템을 장악하거나 접근하면 안되는 정보에 접근할 수 있다.

Direct Execution

- 위에서 언급된 Process execution을 수행하는 하나의 방법
- 그냥 CPU 위에서 Direct하게 돌려! - Program 실행 중간에 별다른 처리 없이
 - OS는 Program에 대한 초기, 후기 작업만을 담당

OS	Program
1. Process List의 entry를 생성	
2. Program에 대한 Memory 할당	
3. Program을 Memory에 로드	

4. Stack Set-up (argc와 argv로)	
5. 레지스터 초기화	
6. Execute call main()	
	7. main() 실행
	8. main() return
9. Process의 memory free	
10. process list에서 제거	



Running programs에 limit을 걸지 않는다면
OS는 제어권이 없는 것이고 그냥 Library로서 동작한다고 볼 수 있음

띠용용?

3번 단계) Program을 Memory에 로드

- 실행 가능한 프로그램의 **Code(코드)**와 **Data(전역 데이터)**를 디스크에서 RAM(메인 메모리)으로 불러옴
- 대상: 실행 파일(.exe, .out 등)에 들어 있는 명령어(기계어 코드)와 전역 변수, 정적 변수 등의 데이터
- 목적: CPU가 프로그램을 실행하려면, 디스크에 있는 코드를 직접 읽을 수 없기 때문에 RAM에 올려두어야 함
- 프로그램 자체를 메모리에 준비하는 과정

4번 단계) Stack Set-up

- 프로그램이 실행될 때 사용할 스택 영역을 메모리에 할당하고 초기화하는 작업
 - **Stack(스택)**은 프로그램 코드와는 별개의 영역
- 대상: 지역 변수, 함수의 매개변수, 함수의 반환 주소 등을 저장할 빈 메모리 공간
- 목적: 함수 호출과 반환 시 필요한 데이터를 임시로 저장하기 위한 공간을 마련
- 특히, 앞서 설명한 main() 함수의 argc와 argv 인자도 이 스택에 초기화
- 프로그램이 실행될 때 사용할 작업 공간을 준비하는 과정

요약

3번: 책(프로그램 코드)을 도서관(메모리)에 가져다 놓는 것

| 4번: 그 책을 읽고 필기하며 계산할 책상(스택)을 준비하는 것

Recap: Process Creation

1. Load

- OS는 Program **code**를 메모리에 Load → process의 address space로
- Program은 원래 disk에 **executable format**으로 저장되어 있음 ex) a.out
- 이때 OS는 **lazy**한 방식으로 process를 load함

Lazy Loading

code와 data 중 program execution에 필수적인 일부분만 Loading하는 것

2. Run-Time Stack Set-Up

- local variables, function parameters, return address등
- main함수의 Parameter인 **argv 배열, argc도 초기화**
- **Register 초기화**도 이 단계에서 진행 - 인터넷에서 찾은 거

3. Heap Creation

- Explicit Dynamic Memory Allocation시 필요한 heap
 - ex) malloc(), free()

4. 추가적인 initialization tasks

- **Input/output (I/O) setup**
 - Process의 기본 open file descriptors 3개: STDIN, STDOUT, STDERR

| 1, 2, 3, 4는 OS가 하는 것이라고 보면 됨, 이후 process에게 CPU 제어권 넘겨주는 것

5. Start the program

- main()에서 시작

- CPU의 제어권을 새로 만들어진 **process**에게 **control** 넘겨줌

Issues

1. User can do wrong thing

```
int *i;
i = 0 ;
*i = 1 ;
```

2. Getting the control back from CPU is not easy

```
i = -1;
while (i < 0)
    do something;
```

Problem 1 : Process가 Restricted Operations를 수행하고자 할 때 문제가 발생

- Disk에 대한 I/O Request
- CPU / Memory와 같은 System Resource에 대한 접근 요청

Solution 1: Grant everything


Solution 2: Use protected control transfer

- **User mode**: Applications은 hardware resources에 대한 full access 불가능
- **Kernel mode**: The OS has access to the full resources of the machine
 - 인터넷에서 찾은거
 - kernel mode는 또 supervisor mode, system mode, privileged mode로 구분
 - 특정 instruction은 privileged일 때만 사용 가능 ex) I/O Request, File System Request

- 즉, User Mode에 있는 Process는 Non-privileged Operation만 수행 가능
 - 그러나!!! User Mode의 Program도 Privileged(Restrictive) Operation이 필요함
 - 따라서 User는 System Call과 같은 Interface를 통해 Kernel Mode 전환

System Call

- Accessing the file system, Creating and destroying processes, Communicating with other processes, Allocating more memory...
- User Mode를 Kernel Mode로 전환, Return할 때, User로 돌아감
 - system call 함수는 내부적으로 Trap (Interrupt)



```
#define syscall3(NUMBER, ARG0, ARG1, ARG2)
({
    int retval;
    asm volatile
    ("pushl %[arg2]; pushl %[arg1]; pushl %[arg0]; "
     "pushl %[number]; int $0x30; addl $16, %%esp"
     : "=a" (retval)
     : [number] "i" (NUMBER),
       [arg0] "g" (ARG0),
       [arg1] "g" (ARG1),
       [arg2] "g" (ARG2)
     : "memory");
    retval;
})
```

arg2, arg1, arg0, syscall number 넣고, interrupt 명령어 걸고 다시 복귀하면 스택 포인터 증가시키기

Register syscall handler at IDT

Interrupt descriptor table

- Table of the locations of the interrupt handlers for the associated interrupt number.
- Register the syscall handler at interrupt 0x30(src/userprog/syscall.c).

```
void
syscall_init (void)
{
    intr_register_int (0x30, 3, INTR_ON, syscall_handler, "syscall");
}
```

Interrupt Descriptor table

0x0	divide_error()
	debug()
	nmi()
	...
0x30	syscall_handler()
	...

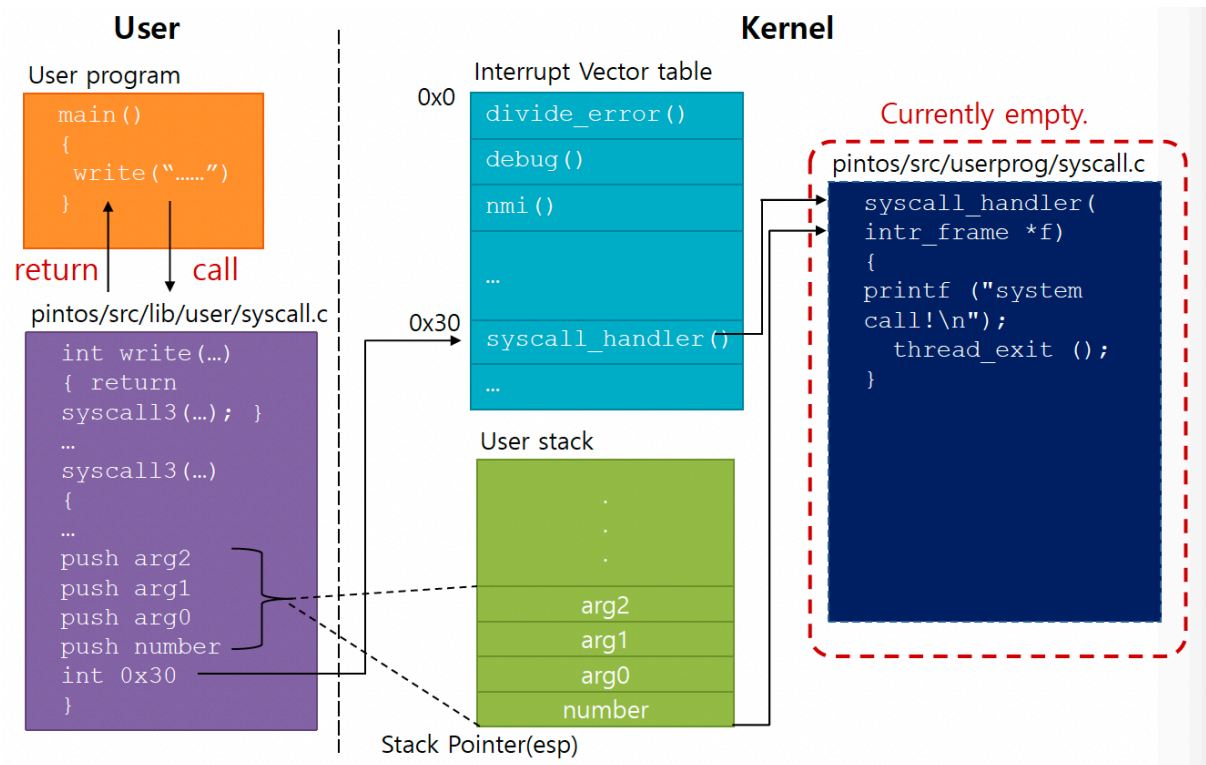
int \$0x30 실행 과정:

1. CPU가 IDT 배열에서 `idt[0x30]` 읽음
2. 거기 저장된 주소(0xC0105A00)를 가져옴
3. 그 주소로 PC 점프
4. 그때부터 실제 명령어 실행

정리:

- IDT = 주소 배열 (데이터)
- 핸들러 = 실제 명령어 (코드)
- IDT는 핸들러를 가리키는 포인터들의 모음!

Call flow(Pintos): write()



kernel stack이랑 process-struct 차이점

cache는 context switch에 포함안된다

Key Concept of System Call

Trap instruction

int n

- Raise the privilege level to kernel mode
- Save registers at the kernel stack. (eip, cs, eflags, esp, ss)
- Locates the destination in the kernel - 이전에 말한 IDT?
- Jumps to the destination.

- Interrupt number for system call
 - 0x64 @xv6 | 0x30 @Pintos | e.g. int 0x64

Return-from-Trap instruction

- Return into the calling user program
- Reduce the privilege level back to user mode

What code to run in trap?

Sources of trap

- Completion of disk IO
- Keyboard interrupt
- System call

Trap handler

- the code to run for each interrupt number (trap number)

Trap table

- The address of the trap handlers.
- The OS informs the hardware of the locations of trap handlers.
- The instruction to inform the location of the trap table is also privileged instruction.
- You cannot execute this instruction in user mode.

Limited Direction Execution

OS @ boot(kernel mode)	Hardware	
initialize trap table		

	remember address of ... syscall handler	
OS @ run(kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg / PC return-from -trap		
	restore regs from kernel stack move to user mode jump to main	
		Run main() ... Call system trap into OS
	save regs to kernel stack move to kernel mode jump to trap handler	
Handle trap Do work of syscall return-from-trap		
	restore regs from kernel stack move to user mode jump to PC after trap	
		return from main trap (via exit())
Free memory of process Remove from process list		

Problem 2: Switching Between Processes

- How can the OS regain control of the CPU so that it can switch between processes?
 - A cooperative Approach: **Wait for system calls**
 - A Non-Cooperative Approach: **The OS takes control**

Solution 1: A cooperative Approach: Wait for system calls

- Processes **periodically give up the CPU** by making **system calls** such as yield.
 - The OS decides to run some other task.
 - Application also transfer control to the OS when they do something illegal.
 - Divide by zero
 - Try to access memory that it shouldn't be able to access
- Ex) Early versions of the Macintosh OS, The old Xerox Alto system

| **A process gets stuck in an infinite loop. → Reboot the machine**

Solution 2: A Non-Cooperative Approach: OS Takes Control

A timer interrupt

- During the boot sequence, the OS start the timer.
- The timer raise an interrupt every few milliseconds.
- When the interrupt is raised
 - The currently running process is suspended.
 - Save enough of the state of the process.
 - A pre-configured interrupt handler in the OS runs

| **A timer interrupt gives OS the ability to run again on a CPU.**

Saving and Restoring Context

- Scheduler makes a decision
 - Whether to continue running the **current process**, or switch to a **different one**.

- If the decision is made to switch, the OS executes context switch

Context Switch

- A low-level piece of assembly code
 - **Save a few register values** for the current process onto its kernel stack
 - General purpose registers
 - PC
 - kernel stack pointer
 - **Restore a few** for the soon-to-be-executing process from its kernel stack
 - **Switch to the kernel stack** for the soon-to-be-executing process

Limited Direct Execution Protocol (Timer interrupt)

OS @ boot (kernel mode)	Hardware	
initialize trap table		
	remember address of ... syscall handler timer handler	
start interrupt timer		
	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A ...
	timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler	
	(Cont.)	

Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) return-from-trap (into B)		
	restore regs(B) from k-stack(B) move to user mode jump to B's PC	
		Process B ...

```

1 # void swtch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax           # put old ptr into eax
9     popl 0(%eax)                # save the old IP
10    movl %esp, 4(%eax)           # and stack
11    movl %ebx, 8(%eax)           # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18    # Load new registers
19    movl 4(%esp), %eax           # put new ptr into eax
20    movl 28(%eax), %ebp          # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp          # stack is switched here
27    pushl 0(%eax)               # return addr put in place
28    ret                          # finally return into new ctxt

```

Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations
 - **Disable interrupts** during interrupt processing
 - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures.